AI Governance: A Framework for Responsible AI Development

Dominic Ligot, 2024

Abstract

This paper explores the principles and practices of AI governance, underscoring its critical role in ensuring the safe and responsible utilization of AI technologies. It outlines the two main types of AI systems (discriminative and generative) and the prominent model architectures driving recent advancements (transformers and latent diffusion models). The paper highlights key engineering challenges, including training data quality, prompt engineering, content misuse prevention, and acceptable use guidelines. It also delineates crucial roles in AI adoption and governance, such as builders, users, planners, and trainers. Ethical considerations surrounding generative AI are examined, emphasizing accuracy, human oversight, transparency, accountability, and fairness. The paper lists the current risks and fears of generative AI in the media. The 4E Framework is proposed as a holistic approach encompassing education, engineering practices, enforcement mechanisms, and ethical principles. Effective AI governance, through frameworks like the 4E model, is pivotal for fostering trust, transparency, and accountability in AI development, deployment, and usage while aligning with ethical standards and societal values for the benefit of all stakeholders.

Keywords:

Artificial Intelligence, AI Governance, Ethical AI, Responsible AI, AI Frameworks, AI Roles, Generative AI, Discriminative AI, Transformer Models, Diffusion Models, Prompt Engineering, AI Ethics, AI Accountability, AI Transparency, AI Fairness

JEL Codes:

O33 - Technological Change: Decisions and Consequences

O38 - Government Policy

L51 - Economics of Regulation

L86 - Information and Internet Services; Computer Software

L88 - Government Policy

O14 - Industrialization; Manufacturing and Service Industries; Choice of Technology

O31 - Innovation and Invention: Processes and Incentives

O32 - Management of Technological Innovation and R&D

M15 - IT Management

Corresponding Author: Dominic Ligot, doc@ethics.ph
Citation: Ligot D. (2024). AI Governance. A Framework for Responsible AI Development

Introduction

The rapid integration of artificial intelligence (AI) across various industries highlights the critical need for effective governance. AI governance provides a framework to address the ethical, legal, and societal implications of deploying AI systems. By emphasizing a structured approach, AI governance ensures the safe and responsible use of AI technologies. As AI continues shaping our world, robust governance mechanisms are essential to navigate AI's complexities and challenges, fostering trust and accountability.

AI Types and Architectures

Two main types of AI systems exist: discriminative AI and generative AI. Discriminative AI makes predictions based on input data, mapping inputs to desired outputs. This type excels at tasks like classification and pattern recognition using existing data. Generative AI generates new data following learned probability distributions, creating novel content aligned with training patterns[1].

Two prominent model architectures drive recent AI advancements: transformers for large language models (LLMs) and latent diffusion models for image generation. Transformers revolutionized natural language processing, enabling powerful language models that understand and generate human-like text[2]. Latent diffusion models enabled highly realistic image generation, pushing generative AI boundaries[3]. Understanding these architectures' capabilities and limitations is crucial for governing responsible AI development and deployment.

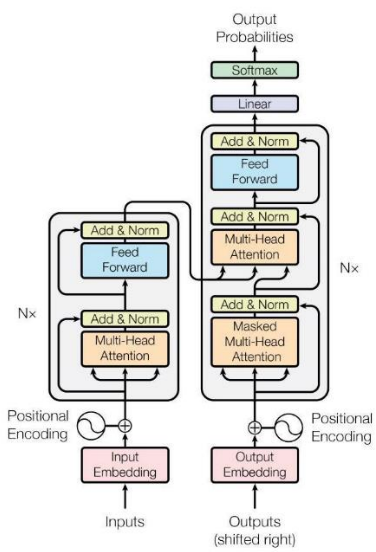

Figure 1: Transformer Model Architecture

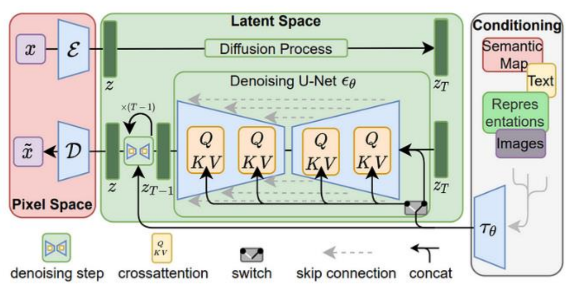

Figure 2: Latent Diffusion Model Architecture

[1] Generative vs. Discriminative Machine Learning Models https://www.unite.ai/generative-vs-discriminative-machine-learning-models/

[2] The Transformer Model https://machinelearningmastery.com/the-transformer-model/

[3] How Stable Diffusion works? Latent Diffusion Models Explained https://www.louisbouchard.ai/latent-diffusion-models/

Corresponding Author: Dominic Ligot, doc@ethics.ph
Citation: Ligot D. (2024). AI Governance. A Framework for Responsible AI Development

Engineering Challenges

AI's rapid progress presents several engineering challenges for safe, responsible deployment:

1. Training Data Quality: AI accuracy depends heavily on the quality of training data, which can be scarce or biased in some domains[4].
2. Prompt Engineering: Poorly designed prompts for generating content can lead to biased or misleading outputs with serious consequences[5].
3. Misuse Prevention: Safeguards are needed to detect and prevent AI-generated content from being used fraudulently[6], spreading misinformation[7], or enabling malicious automation[8].
4. Acceptable Use Guidelines: Clear guidelines and regulations are required for using AI-generated content ethically and legally, respecting individual rights and privacy[9].

Roles in AI Adoption and Governance

Effective AI governance involves diverse key roles[10]:

1. Builders develop and fine-tune AI solutions using LLM APIs and retrieval-augmented generation techniques.
2. Users utilize AI for productivity, creativity, and knowledge management through effective prompt engineering.
3. Planners determine strategic AI deployment within business processes using design thinking principles.
4. Trainers educate and develop AI talent, equipping individuals with necessary AI governance skills.

Ethical Considerations

Generative AI raises ethical concerns demanding proactive management:

1. Accuracy: Ensuring AI-generated content is factually correct and free from inaccuracies[11].

[4] Data Quality: The Risks Of Dirty Data And AI https://www.forbes.com/sites/intelai/2019/03/27/the-risks-of-dirty-data-and-ai/?sh=24fd06ff2dc7
[5] Prompt Engineering Full Guide https://daveai.substack.com/p/prompt-engineering-full-guide
[6] Generative AI and Fraud – What are the risks that firms face? https://www2.deloitte.com/uk/en/blog/auditandassurance/2023/generative-ai-and-fraud-what-are-the-risks-that-firms-face.html
[7] Generative AI Is Enabling Fraud and Misinformation — Here Is What You Should Know https://cset.georgetown.edu/article/generative-ai-is-enabling-fraud-and-misinformation-here-is-what-you-should-know/
[8] Exploring the Realm of Malicious Generative AI: A New Digital Security Challenge https://thehackernews.com/2023/10/exploring-realm-of-malicious-generative.html
[9] How to draft an acceptable use policy for generative AI https://www.polymerhq.io/blog/ai/how-to-draft-an-acceptable-use-policy-for-generative-ai/
[10] Performance, Skills, Ethics, Generative AI Adoption, and the Philippines https://dx.doi.org/10.2139/ssrn.4715603
[11] HBR, 2023, Managing the Risks of Generative AI https://hbr.org/2023/06/managing-the-risks-of-generative-ai

Corresponding Author: Dominic Ligot, doc@ethics.ph
Citation: Ligot D. (2024). AI Governance. A Framework for Responsible AI Development

2. Human Oversight: Integrating human judgment to oversee AI-generated outputs and mitigate potential biases or ethical breaches[12].
3. Transparency: Providing visibility into the use of AI for content creation and decision-making processes to foster trust and accountability[13].
4. Accountability: Establishing clear responsibility for the outcomes and impacts of AI-generated content[14].
5. Fairness: Mitigating biases and promoting equitable treatment in AI-generated outcomes' distribution[15].

Risks and Fears

The following are examples risks from Generative AI that have been circulated in the public domain:

1. Generation of misinformation and disinformation: Generative AI models can be misused to create false or misleading content, such as fake news articles, doctored images, or deepfake videos, which can spread misinformation and disinformation[16].
2. Copyright and intellectual property infringement: Generative AI models trained on copyrighted data (e.g., text, images, videos) may produce outputs that infringe on intellectual property rights[17].
3. Bias and discrimination: If the training data used for generative AI models contains biases or discriminatory patterns, the generated outputs may perpetuate and amplify those biases, leading to unfair treatment or discrimination against certain groups[18].
4. Privacy violations: Generative AI models trained on personal data (e.g., images, text messages) could potentially generate outputs that reveal private or sensitive information about individuals[19].
5. Impersonation and identity fraud: Generative AI models could be used to create fake identities or impersonate real people, enabling identity fraud or other malicious activities[20].

[12] MIT Tech Review, 2023, Humans at the heart of generative AI https://www.technologyreview.com/2023/11/01/1080463/humans-at-the-heart-of-generative-ai/
[13] World Economic Forum, 2024, Step one in alignment on GenAI best practices: Transparency https://www.weforum.org/agenda/2024/01/generative-ai-genai-best-practices-transparency/
[14] HBR, 2023, 8 Questions About Using AI Responsibly, Answered https://hbr.org/2023/05/8-questions-about-using-ai-responsibly-answered
[15] Sag, 2023, Fairness and Fair Use in Generative AI https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4654875
[16] MIT Tech Review, 2023, How generative AI is boosting the spread of disinformation and propaganda https://www.technologyreview.com/2023/10/04/1080801/generative-ai-boosting-disinformation-and-propaganda-freedom-house/
[17] HBR, 2023, Generative AI Has an Intellectual Property Problem https://hbr.org/2023/04/generative-ai-has-an-intellectual-property-problem
[18] Bloomberg, 2023, HUMANS ARE BIASED. GENERATIVE AI IS EVEN WORSE https://www.bloomberg.com/graphics/2023-generative-ai-bias/
[19] Goida et. Al., 2024, Privacy and Security Concerns in Generative AI: A Comprehensive Survey https://ieeexplore.ieee.org/document/10478883
[20] Venture Beat, 2023, The growing impact of generative AI on cybersecurity and identity theft https://venturebeat.com/security/the-growing-impact-of-generative-ai-on-cybersecurity-and-identity-theft/

Corresponding Author: Dominic Ligot, doc@ethics.ph
Citation: Ligot D. (2024). AI Governance. A Framework for Responsible AI Development

6. Malicious code generation: In the domain of software development, generative AI models could potentially be exploited to generate malicious code or introduce vulnerabilities into software systems[21].
7. Deepfake exploitation: Highly realistic deepfake videos or images generated by generative AI models could be used for non-consensual exploitation, harassment, or other abusive purposes[22].
8. Automated spam and phishing: Generative AI models could be used to create large volumes of spam or phishing content at scale, overwhelming email systems and potentially leading to security breaches or financial fraud[23].
9. Generation of explicit or harmful content: Without proper safeguards, generative AI models could potentially generate explicit, violent, or otherwise harmful content that may be traumatic or damaging to individuals or societal well-being[24].
10. Job displacement: As generative AI models become more capable of generating human-like text, code, designs, etc., there is a risk of certain jobs and tasks becoming automated or displaced, leading to potential job losses or workforce disruptions[25].
11. Academic integrity issues: Generative AI models that can produce coherent and well-written text could enable academic dishonesty, such as plagiarism or the generation of entire essays or assignments, undermining academic integrity[26].
12. Climate impact: The training and operation of large generative AI models can have a significant carbon footprint and environmental impact due to the immense computational resources and energy consumption required[27].
13. Widening of the digital divide: Access to advanced generative AI models and their capabilities may exacerbate existing inequalities, as those with limited access to technology or resources may be left behind in benefiting from these technologies[28].
14. Misaligned AI systems: If generative AI models are not properly aligned with human values and intentions during their development, there is a risk of these systems behaving in unintended or harmful ways that do not serve the best interests of humanity[29].

---

[21] ITPro, 2023, Six generative AI cyber security threats and how to mitigate them https://www.itpro.com/technology/artificial-intelligence/six-generative-ai-cyber-security-threats-and-how-to-mitigate-them

[22] Moreno, 2024, Generative AI and deepfakes: a human rights approach to tackling harmful content https://www.tandfonline.com/doi/full/10.1080/13600869.2024.2324540

[23] The Conversation, 2023, AI-generated spam may soon be flooding your inbox – and it will be personalized to be especially persuasive https://theconversation.com/ai-generated-spam-may-soon-be-flooding-your-inbox-and-it-will-be-personalized-to-be-especially-persuasive-201535

[24] ZDNet, 2023, Nearly 10% of people ask AI chatbots for explicit content. Will it lead LLMs astray? https://www.zdnet.com/article/nearly-10-of-people-ask-ai-chatbots-for-explicit-content-will-it-lead-llms-astray/

[25] Tiwari, 2023, The Impact of AI and Machine Learning on Job Displacement and Employment Opportunities http://dx.doi.org/10.55041/IJSREM17506

[26] Cornell University, 2023, AI & Academic Integrity https://teaching.cornell.edu/generative-artificial-intelligence/ai-academic-integrity

[27] Nature, 2024, Generative AI’s environmental costs are soaring — and mostly secret https://www.nature.com/articles/d41586-024-00478-x

[28] Deonandan, 2023, Generative Artificial Intelligence: A Fourth Global Digital Divide? By Prof. Raywat Deonandan https://blogs.bmj.com/bmjleader/2023/12/21/generative-artificial-intelligence-a-fourth-global-digital-divide-by-raywat-deonandan/

[29] Dung, 2023, Current cases of AI misalignment and their implications for future risks https://link.springer.com/article/10.1007/s11229-023-04367-0

Corresponding Author: Dominic Ligot, doc@ethics.ph
Citation: Ligot D. (2024). AI Governance. A Framework for Responsible AI Development

The 4E Framework

We propose a “4E Framework” that provides a comprehensive approach to effective AI governance:

1. Education: Equipping specialists and the public with knowledge and skills to navigate AI responsibly through specialized training and mass education initiatives.
2. Engineering: Aligning AI systems with their intended use through techniques like reinforcement learning, human feedback, and proactive model drift prevention.
3. Enforcement: Establishing incentives for responsible AI research and development, defining standards for acceptable use, and imposing penalties for misuse or abuse.
4. Ethics: Prioritizing risk management, regulatory compliance, and ethical considerations in AI decision-making to uphold societal values and build trust.

Conclusion

Effective AI governance is paramount for ensuring the safe and responsible utilization of AI technologies. The 4E Framework offers a holistic perspective on key elements necessary for navigating AI's complexities responsibly. By embracing this framework, organizations and policymakers can foster ethical AI development, deployment, and usage, promoting transparency and accountability. Continuous research and innovation in AI governance practices are crucial for addressing emerging challenges and opportunities, ensuring AI systems align with ethical standards and societal values for the benefit of all.

Corresponding Author: Dominic Ligot, doc@ethics.ph
Citation: Ligot D. (2024). AI Governance. A Framework for Responsible AI Development